

# AI Meeting Assistant SaaS

Project Brief & Build Plan — prepared for continuation with Claude

**Note to Claude (any future session):** You are acting as a genius SaaS builder — a senior product architect and hands-on engineering partner (via Claude Code) for a solo Indian founder building an AI meeting-notes SaaS in the style of Quill Meetings. This document is the full context and decision record from prior planning conversations. Read it fully before responding. Treat every decision marked “LOCKED” below as final unless the founder explicitly reopens it. Continue execution from the phase the founder indicates, ask only what is still genuinely open, and keep giving honest, realistic estimates rather than optimistic ones.

## Decisions locked in so far

| Area                     | Decision                                                                                                                                                                                                                               |
|--------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Product                  | AI meeting assistant SaaS — records/joins calls, transcribes, generates summaries, action items, and decisions (Quill Meetings-style).                                                                                                 |
| Meeting platforms        | <b>Zoom + Google Meet only</b> for v1. Microsoft Teams explicitly deferred (see Open Risks below) — do not build it unless the founder reopens this.                                                                                   |
| Target feature ceiling   | “Full enterprise” tier: SSO/SAML, admin controls, custom AI personalization per customer, white-label. This is the end-state across all 4 phases, not the v1 scope.                                                                    |
| Who builds it            | Solo founder, coding hands-on with Claude Code as pair-programmer. No hired dev team, no outsourced agency.                                                                                                                            |
| Working capacity assumed | ~6–7 focused hours/day, 6 days/week (~38–40 hrs/week), full-time on this project. <b>If actual availability is part-time (3–4 hrs/day), every phase below roughly doubles in duration.</b> Confirm this before trusting the deadlines. |
| Phase 1 deadline         | 6–8 weeks for a real, sellable MVP (not the full enterprise scope — that was an explicit correction made during planning; see Roadmap).                                                                                                |
| Currency / FX basis      | All ■ figures computed at ~■96/US\$1 (mid-July 2026 mid-market rate). Re-check the live rate before treating these as final if significant time has passed.                                                                            |

## Roadmap & Phases

Reality check baked into this plan: attempting the full enterprise scope solo in 6–8 weeks does not fit together — SSO/SAML, admin console, custom AI personalization, and white-label on top of a working meeting bot is realistically a 6–9 month build even with heavy Claude Code acceleration. The plan below sequences it so Phase 1 is genuinely shippable in the founder's target window, with enterprise features layered afterward.

| Phase | Duration | What ships |
|-------|----------|------------|
|-------|----------|------------|

|                                              |                                                         |                                                                                                                                                                                                                                |
|----------------------------------------------|---------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Phase 1<br/>Core MVP</b>                  | Weeks 1–8                                               | Bot joins Google Meet + Zoom · live transcription · Claude-generated summary/action items/decisions · meeting dashboard · keyword search · Google Calendar auto-join · Stripe billing · basic email + Google auth (no SSO yet) |
| <b>Phase 2<br/>Team layer</b>                | Months 3–4 (~5–6 wks, shortened since Teams is dropped) | Slack + HubSpot/Salesforce push · team workspaces & roles · semantic/vector search · usage analytics dashboard                                                                                                                 |
| <b>Phase 3<br/>Enterprise hardening</b>      | Months 5–6                                              | SSO/SAML + SCIM provisioning (via WorkOS or equivalent) · admin console & audit logs · white-label (custom domain/branding) · security review · DPA/compliance documentation                                                   |
| <b>Phase 4<br/>Scale &amp; customization</b> | Months 7–9                                              | Per-customer AI customization via RAG + tailored system prompts (not literal fine-tuning — see Key Caveats) · load testing/autoscaling · SOC 2 audit engagement, <i>only if</i> enterprise deals require it                    |

**Total realistic timeline: ~9 months** to true full-enterprise readiness, assuming full-time solo work. The 6–8 week deadline applies to Phase 1 only.

## Phase 1 Effort Assumption (why 8 weeks)

Phase 1's 8-week estimate assumes ~300–320 total hours (6–7 hrs/day × 6 days/week × 8 weeks), full-time. Breakdown:

| Workstream                                                             | Est. hours |
|------------------------------------------------------------------------|------------|
| Architecture, infra setup, tech-stack scaffolding                      | 30–40      |
| Meeting-bot integration (Zoom + Meet) — hardest, flakiest part         | 70–90      |
| Transcription pipeline + Claude summarization logic                    | 40–50      |
| Dashboard, auth, calendar integration                                  | 50–60      |
| Stripe billing                                                         | 15–20      |
| Testing, bug fixing, edge cases (dropped audio, bot disconnects, etc.) | 50–60      |

**Critical dependency:** if actual daily availability is part-time (3–4 hrs/day around other work), Phase 1 stretches to roughly 4–5 months. The bot-integration workstream in particular does not compress well with more hours in one sitting — debugging why a Zoom bot silently drops audio at minute 22 takes the time it takes. **Confirm current availability before committing to a date.**

## Cost Estimates (₹, solo founder + Claude Code)

FX basis: ~₹96 per US\$1 (mid-market, mid-July 2026).

| Phase                         | Claude Code (build) | Infra / tools                    | One-off                                    | Phase total |
|-------------------------------|---------------------|----------------------------------|--------------------------------------------|-------------|
| Phase 1 (2 mo)                | ₹45,000             | ₹80,000                          | ₹40,000 (registration, legal-lite, domain) | ₹1.6–2L     |
| Phase 2 (~1.5 mo, Teams-free) | ₹30,000             | ₹90,000                          | —                                          | ₹1.2–1.3L   |
| Phase 3 (2 mo)                | ₹40,000             | ₹2,00,000 (incl. WorkOS for SSO) | ₹75,000 (legal/DPA, security review)       | ₹3.1–3.5L   |
| Phase 4 (3 mo)                | ₹60,000             | ₹3,50,000                        | —                                          | ₹4–4.5L     |

**Total build cost, 9 months, Zoom+Meet scope: roughly ₹10–13 lakh.** Separately, only if/when an enterprise deal specifically requires it: SOC 2 Type 1 audit adds ₹10–25 lakh (auditor + Vanta/Drata-style tooling) — deliberately excluded from the core build budget above.

# Claude Costs — Build vs. Runtime (kept separate deliberately)

## 1. Claude Code — cost to build the product

Claude Max plan runs \$100 or \$200/month and primarily adds usage headroom (20x vs Pro at the \$200 tier) rather than new features, and includes access to Opus 4.8 and Sonnet 5. For solo, heavy daily agentic coding, budget the \$200/month tier — ~**■19,200/month**, with an extra **■5,000–10,000/month** buffer for any API overflow. This is already folded into the phase totals above.

## 2. Claude API — cost to run the product for customers (usage-based, scales with volume)

Published rates used: Sonnet 5 at introductory \$2 input / \$10 output per million tokens through August 31, 2026, moving to standard \$3/\$15 after. Per-meeting math below uses the standard (post-introductory) rate to stay conservative:

| Item                                      | Estimate                                                                              |
|-------------------------------------------|---------------------------------------------------------------------------------------|
| Avg. 45-min meeting transcript            | ~10,000 input tokens                                                                  |
| Summary + action items + decisions output | ~1,000 output tokens                                                                  |
| Cost per meeting                          | $(10,000/1M \times \$3) + (1,000/1M \times \$15) \approx \$0.045 \approx \text{■4.3}$ |

| Meetings / month | Approx. Claude API cost / month |
|------------------|---------------------------------|
| 1,000            | ~ <b>■4,300</b>                 |
| 10,000           | ~ <b>■43,000</b>                |
| 1,00,000         | ~ <b>■4.3 lakh</b>              |

This is a per-customer-usage cost, not a build cost — it grows with revenue, which is favorable unit economics for the SaaS.

## Key Caveats & Open Risks

- **No public Claude fine-tuning:** Anthropic does not offer self-serve fine-tuning of Claude models. “Custom AI per customer” in Phase 4 means retrieval-augmented prompting (per-customer knowledge base + tailored system prompts), not literal fine-tuned weights. Genuine fine-tuning would require a direct conversation with Anthropic's enterprise sales team.
- **Microsoft Teams deferred, not eliminated:** a meaningful share of large enterprise accounts (banks, large IT-services firms, government-adjacent orgs) run primarily on Teams. Skipping it doesn't hurt the SMB/startup customers targeted in Phases 1–3, but if a big account requests it later, plan for a Phase 5. Recommendation carried into the build: pick a bot-infrastructure provider (e.g. a Recall.ai-style service) that is Teams-capable under the hood, so adding it later is closer to a config change than a rebuild.
- **SOC 2 is a sales unlock, not a launch requirement:** budget **■10–25 lakh** for it only once enterprise prospects are actually asking for it — don't front-load this cost into the MVP budget.
- **All costs are estimates, not quotes:** infra pricing (hosting, bot-infrastructure providers, WorkOS, vector DB), FX rates, and Claude API pricing can all move. Re-verify current numbers before finalizing a budget, especially if this document is being reused more than a few weeks after ~19 July 2026.
- **Tech stack not yet chosen:** leaning toward Next.js + Postgres + a managed meeting-bot-infrastructure provider (e.g. Recall.ai) rather than building a headless-browser bot from scratch, which alone could cost 4–6 weeks of engineering time. This was proposed but not confirmed as of the last conversation.

## Open Questions for the Next Session

---

1. Confirm real daily availability (full-time vs. part-time) — this changes every deadline in this document.
2. Confirm tech stack: framework, database, and specifically which meeting-bot-infrastructure provider to use for Zoom + Google Meet.
3. Confirm transcription provider (e.g. Deepgram vs. AssemblyAI) if not bundled with the chosen bot-infrastructure provider.
4. Re-check current Claude API pricing, USD/INR rate, and third-party tool pricing (WorkOS, hosting, vector DB) if this document is more than a few weeks old.
5. Once confirmed, begin Phase 1, Week 1: architecture + infra scaffolding + first working Zoom or Google Meet bot in a dev environment.

---

Generated from a planning conversation on 19 July 2026. Treat all figures as estimates for planning purposes, not fixed quotes.